

Technical Interview Questions & Scaling Solutions

Potential Interview Questions

Code Structure & Design Questions

1. "Why did you choose a class-based approach instead of just functions?"
 - **Answer:** Encapsulation of state (cache, configuration), reusability, easier testing, and following OOP principles. The class maintains configuration and cache state across multiple operations.
2. "What's the difference between `__init__` parameters and hard-coded values?"
 - **Answer:** Configurability for different environments (dev/staging/prod), testability with mock values, and flexibility without code changes.
3. "Why use a private method `_fetch_lei_data`?"
 - **Answer:** Internal implementation detail that shouldn't be called directly by users. It maintains the class's interface contract and allows internal changes without breaking external code.
4. "Explain the `-> Dict[str, Any]` syntax"
 - **Answer:** Type hints that specify return type. `Dict[str, Any]` means dictionary with string keys and values of any type. Helps with IDE autocomplete, static analysis, and documentation.

Error Handling Questions

5. "Why create a custom exception `LEIEnrichmentError`?"
 - **Answer:** Specific error handling, clearer debugging, allows catching specific errors vs generic ones, better logging and monitoring.
6. "What happens if the API returns malformed JSON?"
 - **Answer:** `json.JSONDecodeError` is caught, logged as error, returns empty result, and caches it to avoid repeated failures.
7. "Why continue processing if one LEI fails?"
 - **Answer:** Resilience - partial success is better than total failure. Business requirement to process as much data as possible.

Performance & Optimization Questions

8. "Why get unique LEIs instead of processing each row?"

- **Answer:** API call optimization. 1000 rows with 50 unique LEIs = 50 calls instead of 1000. Significant time and rate limit savings.

9. "Explain the caching strategy. What are its limitations?"

- **Answer:** In-memory dictionary for session persistence, file-based for cross-session. Limitations: memory usage grows, no cache expiration, single-machine only.

10. "What's exponential backoff and why use it?"

- **Answer:** `2^attempt` creates increasing delays (1s, 2s, 4s...). Prevents overwhelming struggling servers, gives them time to recover.

Data Processing Questions

11. "Why use `df.copy()` instead of modifying the original?"

- **Answer:** Immutability principle, prevents side effects, allows original data to be used elsewhere, safer for debugging.

12. "Explain the `axis=1` parameter in `apply()`"

- **Answer:** `axis=0` applies function to columns, `axis=1` applies to rows. We need row-wise calculation for transaction costs.

13. "Why use `map()` instead of `apply()` for LEI lookups?"

- **Answer:** `map()` is faster for simple lookups, `apply()` is for complex functions. `map()` is optimized for dictionary/function mapping.

Business Logic Questions

14. "Walk me through the transaction cost calculation"

- **Answer:** Country-specific formulas: GB uses interest calculation, NL uses inverse rate with absolute value, others default to 0.

15. "How would you handle new country requirements?"

- **Answer:** Add new elif conditions, externalize rules to config file/database, or use strategy pattern for complex rules.

API Integration Questions

16. "Why use `raise_for_status()`?"

- **Answer:** Automatically converts HTTP error codes (4xx, 5xx) to Python exceptions for consistent error handling.

17. "How do you handle rate limits?"

- **Answer:** Sleep delay between calls, exponential backoff on failures, respect API documentation limits.

18. "What if the API changes its response format?"

- **Answer:** Defensive programming with `.get()` methods, graceful degradation, version the API calls, comprehensive logging.

Scalability & Architecture Questions

19. "How would you test this code?"

- **Answer:** Unit tests with mocked API calls, integration tests with test data, performance tests with large datasets.

20. "How would you monitor this in production?"

- **Answer:** Structured logging, metrics (success rate, processing time), alerts on failures, cache hit rates.

21. "What would you do if processing 1 million records?"

- **Answer:** Batch processing, async I/O, database caching, distributed processing, progress tracking.

22. "How would you handle API key management?"

- **Answer:** Environment variables, secret management systems (AWS Secrets Manager, HashiCorp Vault), never hardcode in code.

Scaling Solutions Implementation

1. Async Processing with `asyncio` and `aiohttp`

python

```
import asyncio
import aiohttp
import pandas as pd
from typing import List, Dict, Any
import time
```

```
class AsyncLEIEnricher:
```

```
    def __init__(self,
        base_url: str = "https://api.gleif.org/api/v1/lei-records",
        max_concurrent: int = 10,
        rate_limit_delay: float = 0.1):
        self.base_url = base_url
        self.max_concurrent = max_concurrent
        self.rate_limit_delay = rate_limit_delay
        self.semaphore = asyncio.Semaphore(max_concurrent)
        self._lei_cache = {}
```

```
    async def _fetch_lei_data_async(self, session: aiohttp.ClientSession, lei_code: str) -> Dict[str, Any]:
```

```
        """Async version of LEI data fetching"""
```

```
        if lei_code in self._lei_cache:
            return self._lei_cache[lei_code]
```

```
    async with self.semaphore: # Limit concurrent requests
```

```
        url = f"{self.base_url}?filter[lei]={lei_code}"
```

```
    try:
```

```
        async with session.get(url, timeout=aiohttp.ClientTimeout(total=30)) as response:
            response.raise_for_status()
            data = await response.json()
```

```
        # Same data extraction logic as before
```

```
        if 'data' in data and len(data['data']) > 0:
```

```
            lei_record = data['data'][0]
            attributes = lei_record.get('attributes', {})
```

```
            legal_name = "
```

```
            entity = attributes.get('entity', {})
```

```
            if entity and 'legalName' in entity:
```

```
                legal_name = entity['legalName'].get('name', "")
```

```
            result = {
```

```
                'legalName': legal_name,
```

```
                'bic': attributes.get('bic', []).get(0) if attributes.get('bic') else "",
```

```
                'country': entity.get('legalAddress', {}).get('country', "") if entity else ""
```

```

    else:
        result = {'legalName': "", 'bic': "", 'country': ""}

        self._lei_cache[lei_code] = result
        await asyncio.sleep(self.rate_limit_delay) # Async sleep
        return result

```

```

except Exception as e:
    print(f"Error fetching {lei_code}: {e}")
    result = {'legalName': "", 'bic': "", 'country': ""}
    self._lei_cache[lei_code] = result
    return result

```

```

async def enrich_dataset_async(self, input_data: pd.DataFrame) -> pd.DataFrame:

```

```

    """Async enrichment of entire dataset"""

```

```

    unique_leis = input_data['lei'].unique()

```

```

    async with aiohttp.ClientSession() as session:

```

```

        # Create tasks for all LEIs

```

```

        tasks = [
            self.fetch_lei_data_async(session, lei_code)
            for lei_code in unique_leis
        ]

```

```

        # Execute all tasks concurrently

```

```

        results = await asyncio.gather(*tasks, return_exceptions=True)

```

```

        # Build lei_info dictionary

```

```

        lei_info = {}
        for lei_code, result in zip(unique_leis, results):
            if isinstance(result, Exception):
                lei_info[lei_code] = {'legalName': "", 'bic': "", 'country': ""}
            else:
                lei_info[lei_code] = result

```

```

        # Apply results to DataFrame (same as before)

```

```

        enriched_data = input_data.copy()
        enriched_data['legalName'] = enriched_data['lei'].map(
            lambda x: lei_info.get(x, {}).get('legalName', "")
        )
        enriched_data['bic'] = enriched_data['lei'].map(
            lambda x: lei_info.get(x, {}).get('bic', "")
        )

```

```

    return enriched_data

```

Usage:

```
async def main_async():
    enricher = AsyncLEIEnricher(max_concurrent=10)
    df = pd.read_csv("sample_input.csv")

    start_time = time.time()
    enriched_df = await enricher.enrich_dataset_async(df)
    end_time = time.time()

    print(f"Async processing took: {end_time - start_time:.2f} seconds")
    return enriched_df
```

Run async code:

```
# enriched_df = asyncio.run(main_async())
```

Key Async Concepts:

- `asyncio.Semaphore(10)`: Limits concurrent connections to prevent overwhelming the API
- `aiohttp.ClientSession`: Reuses connections for efficiency
- `asyncio.gather()`: Runs all tasks concurrently and waits for completion
- `async with`: Ensures proper resource cleanup

2. Database Caching with Redis

python

```
import redis
import json
import pickle
from datetime import datetime, timedelta

class DatabaseCachedLEIEnricher:
    def __init__(self,
        base_url: str = "https://api.gleif.org/api/v1/lei-records",
        redis_host: str = "localhost",
        redis_port: int = 6379,
        cache_ttl: int = 86400): # 24 hours
        self.base_url = base_url
        self.cache_ttl = cache_ttl

        # Redis connection with connection pooling
        self.redis_client = redis.ConnectionPool(
            host=redis_host,
            port=redis_port,
            decode_responses=True,
            max_connections=10
        )
        self.redis = redis.Redis(connection_pool=self.redis_client)

    def _get_cache_key(self, lei_code: str) -> str:
        """Generate Redis cache key"""
        return f"lei:v1:{lei_code}"

    def _fetch_from_cache(self, lei_code: str) -> Dict[str, Any]:
        """Fetch LEI data from Redis cache"""
        cache_key = self._get_cache_key(lei_code)
        cached_data = self.redis.get(cache_key)

        if cached_data:
            try:
                return json.loads(cached_data)
            except json.JSONDecodeError:
                # Handle corrupted cache data
                self.redis.delete(cache_key)
                return None
        return None

    def _store_in_cache(self, lei_code: str, data: Dict[str, Any]):
        """Store LEI data in Redis cache with TTL"""
        cache_key = self._get_cache_key(lei_code)
```

```

cache_data = {
    'data': data,
    'cached_at': datetime.utcnow().isoformat(),
    'version': '1.0'
}

try:
    self.redis.setex(
        cache_key,
        self.cache_ttl,
        json.dumps(cache_data)
    )
except Exception as e:
    print(f"Failed to cache data for {lei_code}: {e}")

def _fetch_lei_data(self, lei_code: str) -> Dict[str, Any]:
    """Enhanced fetch with database caching"""
    # Try cache first
    cached_result = self._fetch_from_cache(lei_code)
    if cached_result:
        return cached_result['data']

    # Fallback to API (same logic as original)
    result = self._fetch_from_api(lei_code)

    # Cache the result
    self._store_in_cache(lei_code, result)

    return result

def get_cache_stats(self) -> Dict[str, int]:
    """Get cache performance statistics"""
    cache_keys = self.redis.keys("lei:v1:*")
    total_keys = len(cache_keys)

    # Get cache size in bytes
    cache_memory = sum(
        self.redis.memory_usage(key) or 0
        for key in cache_keys
    )

    return {
        'total_cached_leis': total_keys,
        'cache_memory_bytes': cache_memory,
        'cache_memory_mb': round(cache_memory / (1024 * 1024), 2)
    }

```



```
def clear_cache(self, pattern: str = "lei:v1.*"):
    """Clear cache by pattern"""
    keys = self.redis.keys(pattern)
    if keys:
        self.redis.delete(*keys)
    return len(keys)
```

Database Caching Benefits:

- **Persistence:** Cache survives application restarts
- **Sharing:** Multiple application instances share cache
- **TTL:** Automatic expiration of stale data
- **Memory Management:** Redis handles memory optimization
- **Scalability:** Can handle millions of cache entries

3. Distributed Processing with Celery

python

```
from celery import Celery, group
import pandas as pd
from typing import List
import numpy as np
```

Celery configuration

```
celery_app = Celery(
    'lei_enrichment',
    broker='redis://localhost:6379/0',
    backend='redis://localhost:6379/0'
)
```

```
@celery_app.task(bind=True, max_retries=3)
```

```
def process_lei_batch(self, lei_codes: List[str] -> Dict[str, Dict[str, Any]]:
```

```
    """Celery task to process a batch of LEI codes"""
```

```
    try:
```

```
        enricher = LEIDataEnricher() # Create instance in worker
```

```
        results = {}
```

```
        for lei_code in lei_codes:
```

```
            try:
```

```
                results[lei_code] = enricher._fetch_lei_data(lei_code)
```

```
            except Exception as e:
```

```
                # Log error but continue processing other LEIs
```

```
                print(f"Failed to process {lei_code}: {e}")
```

```
                results[lei_code] = {'legalName': '', 'bic': '', 'country': ''}
```

```
        return results
```

```
    except Exception as exc:
```

```
        # Retry logic for entire batch
```

```
        print(f"Task failed, retrying: {exc}")
```

```
        raise self.retry(countdown=60 * (self.request.retries + 1))
```

```
class DistributedLEIEnricher:
```

```
    def __init__(self, batch_size: int = 50):
```

```
        self.batch_size = batch_size
```

```
    def create_batches(self, lei_codes: List[str] -> List[List[str]]:
```

```
        """Split LEI codes into processing batches"""
```

```
        return [
```

```
            lei_codes[i:i + self.batch_size]
```

```
            for i in range(0, len(lei_codes), self.batch_size)
```

```
        ]
```

```
def enrich_dataset_distributed(self, input_data: pd.DataFrame) -> pd.DataFrame:
```

```
    """Distribute LEI enrichment across multiple workers"""
```

```
    unique_leis = input_data['lei'].unique().tolist()
```

```
    batches = self.create_batches(unique_leis)
```

```
    print(f"Processing {len(unique_leis)} LEIs in {len(batches)} batches")
```

```
    # Create Celery group for parallel execution
```

```
    job = group(process_lei_batch.s(batch) for batch in batches)
```

```
    # Execute all batches in parallel
```

```
    result = job.apply_async()
```

```
    # Wait for all batches to complete
```

```
    batch_results = result.get(timeout=300) # 5 minute timeout
```

```
    # Combine results from all batches
```

```
    lei_info = {}
```

```
    for batch_result in batch_results:
```

```
        lei_info.update(batch_result)
```

```
    # Apply to DataFrame (same as before)
```

```
    enriched_data = input_data.copy()
```

```
    enriched_data['legalName'] = enriched_data['lei'].map(
```

```
        lambda x: lei_info.get(x, {}).get('legalName', '')
```

```
    )
```

```
    enriched_data['bic'] = enriched_data['lei'].map(
```

```
        lambda x: lei_info.get(x, {}).get('bic', '')
```

```
    )
```

```
    return enriched_data
```

```
# Usage:
```

```
def run_distributed_processing():
```

```
    enricher = DistributedLEIEnricher(batch_size=100)
```

```
    df = pd.read_csv("large_input.csv")
```

```
    enriched_df = enricher.enrich_dataset_distributed(df)
```

```
    enriched_df.to_csv("distributed_output.csv", index=False)
```

```
# Start Celery worker:
```

```
# celery -A lei_enricher worker --loglevel=info --concurrency=4
```

Distributed Processing Benefits:

- **Horizontal Scaling:** Add more worker machines
- **Fault Tolerance:** Failed tasks can be retried
- **Load Distribution:** Work spreads across available resources
- **Monitoring:** Built-in task monitoring and statistics

4. Complete Production Architecture

python

```
import asyncio
import aioredis
from celery import Celery
import structlog
from prometheus_client import Counter, Histogram, start_http_server
```

```
class ProductionLEIEnricher:
```

```
    def __init__(self):
```

```
        # Structured logging
```

```
        self.logger = structlog.get_logger()
```

```
        # Metrics
```

```
        self.api_calls = Counter('lei_api_calls_total', 'Total API calls', ['status'])
```

```
        self.processing_time = Histogram('lei_processing_seconds', 'Processing time')
```

```
        # Start metrics server
```

```
        start_http_server(8000)
```

```
    async def enrich_with_monitoring(self, input_data: pd.DataFrame) -> pd.DataFrame:
```

```
        """Production enrichment with full monitoring"""
```

```
        with self.processing_time.time():
```

```
            try:
```

```
                # Log structured data
```

```
                self.logger.info(
```

```
                    "starting_enrichment",
```

```
                    record_count=len(input_data),
```

```
                    unique_leis=len(input_data['lei'].unique())
```

```
                )
```

```
                # Process data
```

```
                result = await self._process_with_circuit_breaker(input_data)
```

```
                # Success metrics
```

```
                self.api_calls.labels(status='success').inc()
```

```
                self.logger.info(
```

```
                    "enrichment_completed",
```

```
                    processed_records=len(result),
```

```
                    success_rate=self._calculate_success_rate(result)
```

```
                )
```

```
            return result
```

```
        except Exception as e:
```

```
self.api_calls.labels(status='error').inc()
self.logger.error("enrichment_failed", error=str(e))
raise
```

```
def _calculate_success_rate(self, df: pd.DataFrame) -> float:
    """Calculate enrichment success rate"""
    total = len(df)
    enriched = len(df[df['legalName'] != ''])
    return round((enriched / total) * 100, 2) if total > 0 else 0
```

Performance Comparison

Method	1000 LEIs	10000 LEIs	Memory Usage	Complexity
Synchronous	~100 seconds	~1000 seconds	Low	Simple
Async (10 concurrent)	~15 seconds	~150 seconds	Medium	Medium
Database Cached	~5 seconds	~50 seconds	Low	Medium
Distributed (4 workers)	~25 seconds	~250 seconds	High	Complex
Combined Approach	~3 seconds	~30 seconds	Medium	High

When to Use Each Approach:

- **Async:** High I/O operations, moderate scale (< 100k records)
- **Database Caching:** Repeated processing, multiple application instances
- **Distributed:** Very large datasets (> 1M records), horizontal scaling needs
- **Combined:** Production systems requiring high performance and reliability